

H₂O₂RAM

A High-Performance Hierarchical Doubly Oblivious RAM

Leqian Zheng¹, Zheng Zhang², Wentao Dong¹, Yao Zhang², Ye Wu², Cong Wang¹

¹City University of Hong Kong, ²ByteDance Inc.

2025 Aug 15

Access patterns matter in encrypted systems

- Access patterns (which encrypted data are accessed) $\xrightarrow{\text{reveal}}$ sensitive information (about encrypted data)

A toy example

```
1 if (secret == true) {  
2     cnt++;  
3 }  
4 else { // do nothing  
5 }
```

Real-world applications

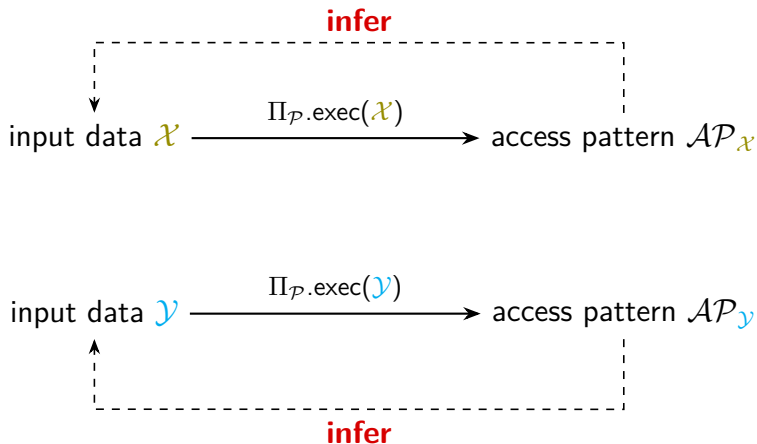
- Searchable encryption
- Private contact discovery for Signal
- Anonymizing Google's Key transparency
- ...

How access patterns pose privacy risks

input data \mathcal{X} $\xrightarrow{\Pi_{\mathcal{P}}.\text{exec}(\mathcal{X})}$ access pattern $\mathcal{AP}_{\mathcal{X}}$

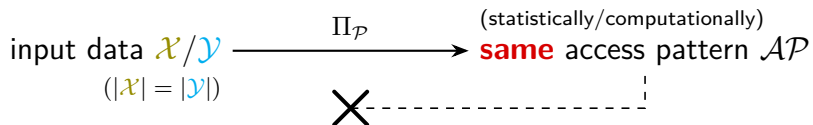
input data \mathcal{Y} $\xrightarrow{\Pi_{\mathcal{P}}.\text{exec}(\mathcal{Y})}$ access pattern $\mathcal{AP}_{\mathcal{Y}}$

How access patterns pose privacy risks



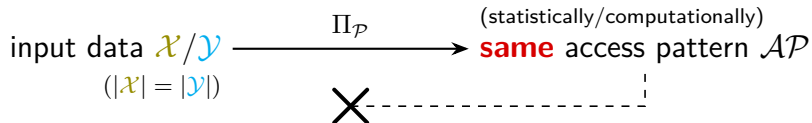
General solution: Oblivious RAM

■ Oblivious RAM (ORAM):



General solution: Oblivious RAM

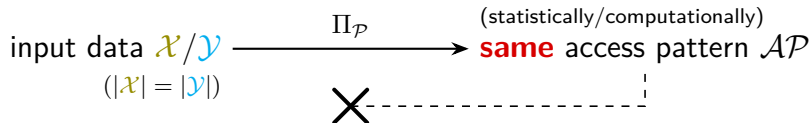
■ Oblivious RAM (ORAM):



- Technique roadmap: tree-based designs (better practical performance) and hierarchical designs (more of theoretical interests)

General solution: Oblivious RAM

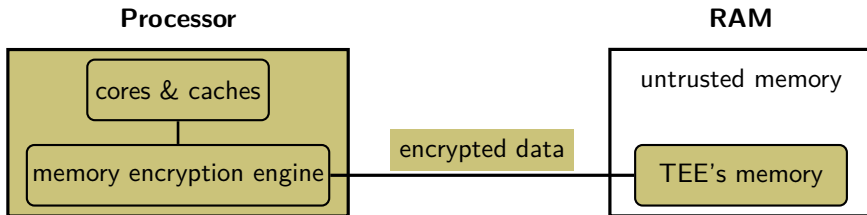
■ Oblivious RAM (ORAM):



- Technique roadmap: tree-based designs (better practical performance) and hierarchical designs (more of theoretical interests)
- The **de facto standard** for concealing access patterns across diverse scenarios and settings

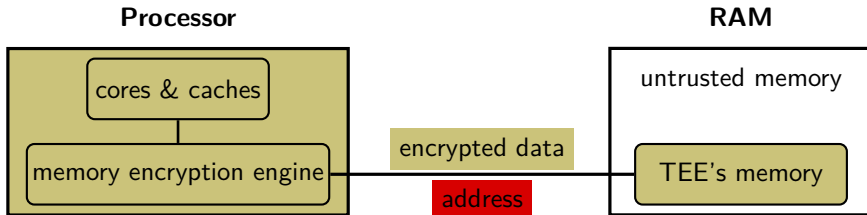
Focused setting: Conceal memory access patterns in TEEs

- Trusted execution environments (TEEs) offer compelling security guarantees with favorable performance and usability 😊



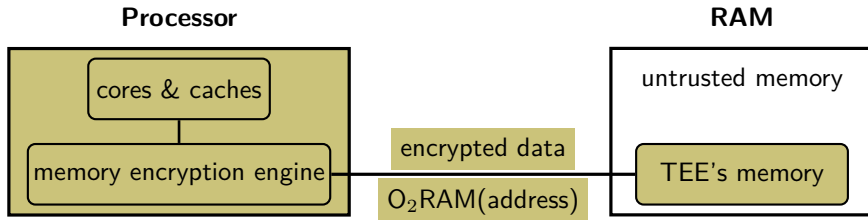
Focused setting: Conceal memory access patterns in TEEs

- Trusted execution environments (TEEs) offer compelling security guarantees with favorable performance and usability 😊
- Memory access patterns remain **exposed by design** in mainstream hardware-based TEEs 😞



Focused setting: Conceal memory access patterns in TEEs

- Trusted execution environments (TEEs) offer compelling security guarantees with favorable performance and usability 😊
- Memory access patterns remain **exposed by design** in mainstream hardware-based TEEs 😞



- General solution for the focused setting: **doubly** oblivious RAM (O₂RAM)

Motivation

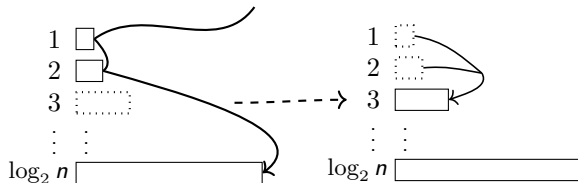
High performance overheads in existing O₂RAM designs:

- Find on a 2^{20} -element map: nanosecs/ 10^{-9} s (plain) \longleftrightarrow millisecs/ 10^{-3} s (oblivious)
- Dijkstra on a graph ($|V| \approx 2^{16}$, $|E| \approx 2^{18}$): millisecs/ 10^{-3} s (plain) $\longleftrightarrow \sim 9$ hours (oblivious)



Amplified by millions — or more!

🤔 $\text{ACCESS}(id) \xrightarrow{\text{handle}}$ (probably amortized) 1000 data blocks in both approaches?

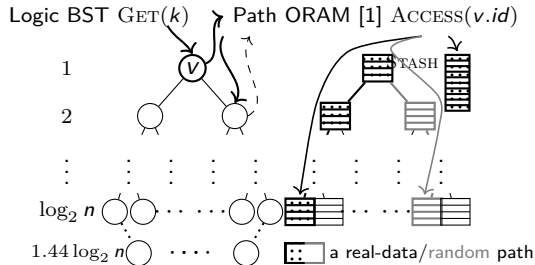


\square : (empty) oblivious hash table

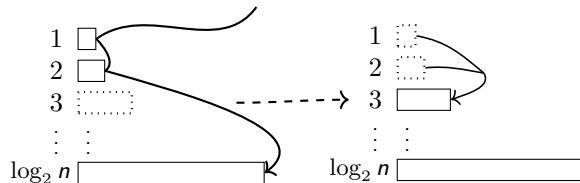
- Better data locality
- Easier to parallelize

Key insight

🤖 Comparable complexity \Rightarrow HORAM performs better!



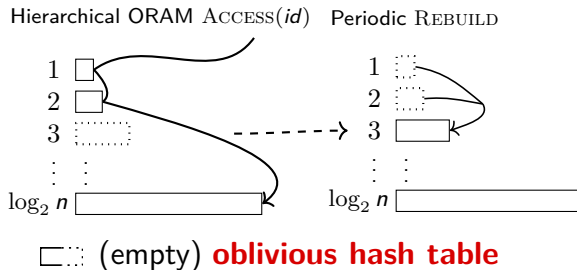
Hierarchical ORAM ACCESS(id) **Periodic** REBUILD



- Better data locality
- Easier to parallelize

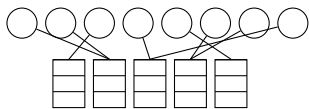
H₂O₂RAM: Technical challenge

- Design high-performance oblivious hash tables

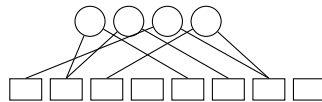


H₂O₂RAM: Our approach

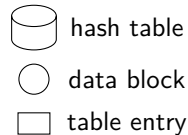
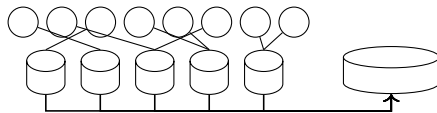
■ Design high-performance oblivious hash tables



(a) Bucket hash



(b) Stashless Cuckoo hash



- (c) Two-tier hash, where a secret number of blocks from the main hash table will be (obviously) relocated to a secondary hash table

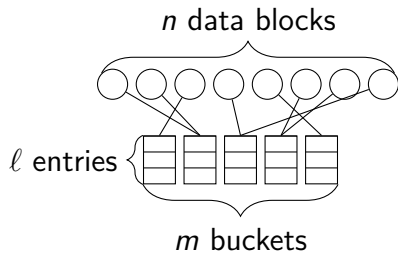
H₂O₂RAM: Core intuitions

- Compute “tight”/appropriate parameters for oblivious hash tables (OHT)

Lifecycle of an OHT (capacity n): 1 **rebuild** + n **lookups**

H₂O₂RAM: Core intuitions

- Compute “tight”/appropriate parameters for oblivious hash tables (OHT)

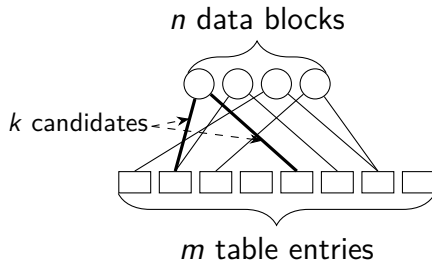


(a) Bucket hash

- each block \rightarrow a bucket via a PRF
- ℓ entries: ensure a negligible probability of bucket overflow
- lookup: linearly scan a bucket (for obliviousness)
- $m \uparrow$: $\ell \downarrow$ lower per-lookup cost, table size $(m \cdot \ell) \uparrow$ higher rebuild cost
- m^* : balance lookup vs. rebuild costs \rightarrow minimal overall cost

H₂O₂RAM: Core intuitions

- Compute “tight”/appropriate parameters for oblivious hash tables (OHT)



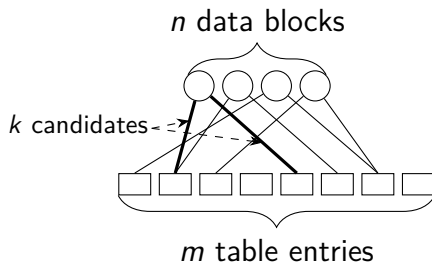
- each block $\rightarrow k$ candidate entries via k PRFs
- $k = 2 w /$ a stash handling unplaceable blocks ($\Omega(\log n)$ for negl overflow prob)

(b) Stashless Cuckoo hash

H₂O₂RAM: Core intuitions

- Compute “tight”/appropriate parameters for oblivious hash tables (OHT)

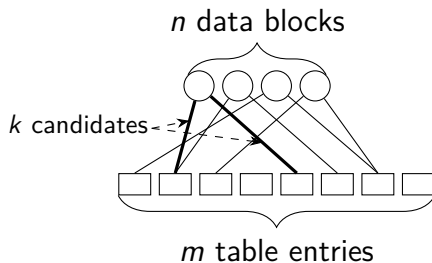
- each block $\rightarrow k$ candidate entries via k PRFs
- $k = ?$ w/o stash + negl overflow prob



(b) Stashless Cuckoo hash

H₂O₂RAM: Core intuitions

- Compute “tight”/appropriate parameters for oblivious hash tables (OHT)

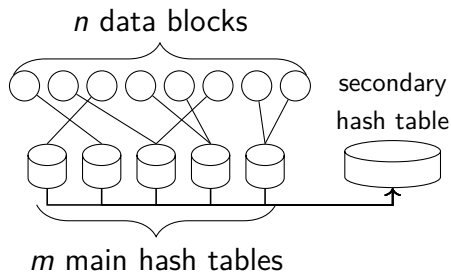


- each block $\rightarrow k$ candidate entries via k PRFs
- $k \in [4, 6]$ suffices (Crypto 23 [3] + our numeric analysis)
- construction: oblivious bipartite matching (potentially of independent interest)

(b) Stashless Cuckoo hash

H₂O₂RAM: Core intuitions

- Compute “tight”/appropriate parameters for oblivious hash tables (OHT)



(c) Two-tier hash

- each block \rightarrow a main HT via a PRF, *nonobliviously*
- ϵ_i (secret) of i^{th} main HT \rightarrow secondary HT, *obliviously*
- to be optimized:
 - #main HTs m
 - sampling parameters for ϵ_i
 - underlying hashing schemes

H₂O₂RAM: Core intuitions

- Compute “tight”/appropriate parameters for oblivious hash tables (OHT)
- Select suitable hash schemes for each H₂O₂RAM level

hash scheme	rebuild time [‡]	lookup time	suitable scenarios
linear scan	$\mathcal{O}(n)$	$\mathcal{O}(n)$	very small n
bucket hash	$\mathcal{O}(n \log^2 n)$	$\mathcal{O}(\ell)$	small to moderate n
stashless Cuckoo hash	$\mathcal{O}(n \log^3 n)$	$\mathcal{O}(1)$	$n < t$, secondary table for two-tier hash
two-tier hash	$\mathcal{O}(n \log^2 \log n)$	$\mathcal{O}(\ell)$	large n

[†] n : hash table size, ℓ : bucket size, t : the number of lookups performed during its lifetime.

[‡] Adopting an $\mathcal{O}(n \log n)$ osort algorithm [2] saves a $\log n$ factor for this column.

H₂O₂RAM: Overall complexity

- Asymptotic comparison of existing O₂RAM schemes, N denotes the capacity

Scheme	Amortized access time
Oblix [4] & GraphOS [5]	$\mathcal{O}(\log^3 N)$
ENIGMAP [6]	$\tilde{O}(\log^2 N)$
H ₂ O ₂ RAM	$\tilde{O}(\log^2 N)$

H₂O₂RAM: Overall complexity

Introduction
○○○Motivation & Key insight
○●H₂O₂RAM
○○○●Evaluations
○

References

References
○

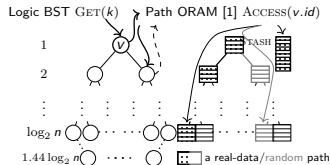
Key insight

■ Asymptoti

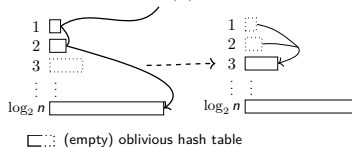


Comparable complexity \Rightarrow HORAM performs better!

≈ capacity



Hierarchical ORAM ACCESS(id) **Periodic** REBUILD



- Better data locality
- Easier to parallelize

H₂O₂RAM

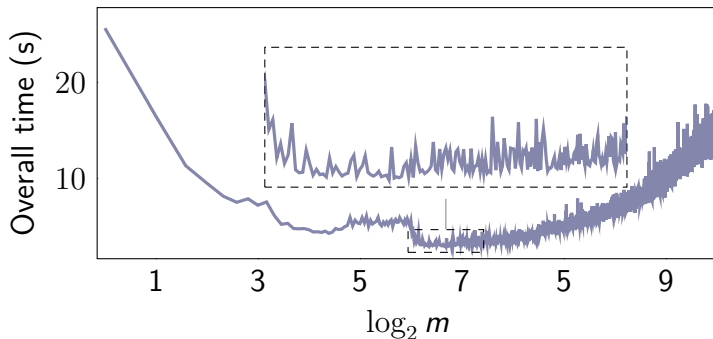
A High-Performance Hierarchical Doubly Oblivious RAM

2025 Aug 15 6 / 12

Comparable complexity achieved!

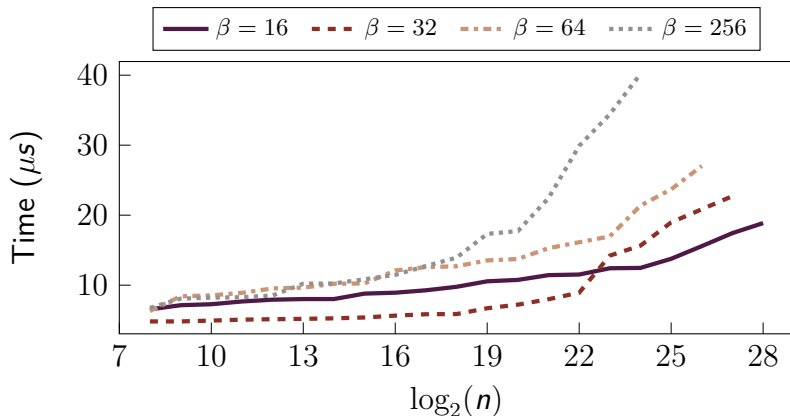
Evaluations: Optimizing OHT parameters

- Overall running time (1 rebuild + n lookups) for a bucket OHT of capacity 8192



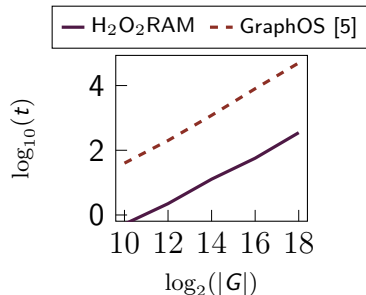
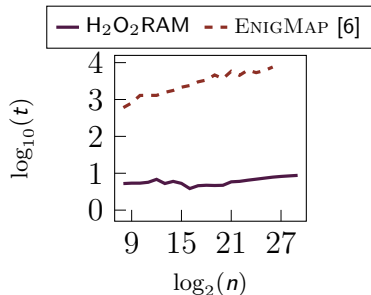
Evaluations: Latency under different settings

- Amortized access time of H₂O₂RAM with different input data sizes n and data block sizes β



Evaluations: Comparative results

- Get operation time t (in microseconds) of a map, where n denotes its capacity
- Single-source shortest path computation time t (in seconds). $|G| = |E| + |V|$ with $|E| = 4|V|$. The results for GraphOS [5] are taken from the “Gr-V0.13E” line in Figure 6c of their paper.



$\sim 10^3$ speedup!

References

- [1] Emil Stefanov et al. “Path ORAM: An extremely simple oblivious RAM protocol”. In: **JACM** 65.4 (2018), pp. 1–26.
- [2] Sajin Sasy, Aaron Johnson, and Ian Goldberg. “Waks-On/Waks-Off: Fast Oblivious Offline/Online Shuffling and Sorting with Waksman Networks”. In: **CCS**. 2023.
- [3] Kevin Yeo. “Cuckoo hashing in cryptography: Optimal parameters, robustness and applications”. In: **CRYPTO**. 2023.
- [4] Pratyush Mishra et al. “Oblix: An efficient oblivious search index”. In: **S&P**. 2018.
- [5] Javad Ghareh Chamani et al. “GraphOS: Towards Oblivious Graph Processing”. In: **VLDB**. 2023.
- [6] Afonso Tinoco, Sixiang Gao, and Elaine Shi. “EnigMap: External-Memory Oblivious Map for Secure Enclaves”. In: **USENIX Security**. 2023.

Thanks for your attention!

Contact:

leqian.zheng@my.cityu.edu.hk